

UL-Inclass-1

March 10, 2020

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from scipy.spatial import distance
from scipy.stats import zscore
```

In class Assignment Expectations/Steps -

Apply Data Cleaning to the Datasets and then apply Kmeans algorithm for find pattern and the best value of the K for the following features.

1. Use features fixed acidity and volatile acidity
2. Use features Citric acidity and fixed acidity
3. Use features residual sugar and sulphates
4. Use features free.sulfur.dioxide and total.sulfur.dioxide
5. Use features fixed acidity, citric acidity and volatile acidity
6. Use features density and pH

```
[2]: df=pd.read_csv('winequality-red.csv')
df.head()
```

```
[2]:
```

	fixed_acidity	volatile_acidity	citric_acid	residual_sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	

	free_sulfur_dioxide	total_sulfur_dioxide	density	pH	sulphates	\
0	11.0	34.0	0.9978	3.51	0.56	
1	25.0	67.0	0.9968	3.20	0.68	
2	15.0	54.0	0.9970	3.26	0.65	
3	17.0	60.0	0.9980	3.16	0.58	
4	11.0	34.0	0.9978	3.51	0.56	

	alcohol	quality
0	9.4	5

1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5

#we have read the data into df and we can see the first five columns of the data clearly which are all numerical

```
[3]: df.shape
```

```
[3]: (1599, 12)
```

#we see the shape of the data as there are 12 columns and 1599 number of rows

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed_acidity          1599 non-null   float64
1   volatile_acidity       1599 non-null   float64
2   citric_acid            1599 non-null   float64
3   residual_sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free_sulfur_dioxide    1599 non-null   float64
6   total_sulfur_dioxide   1599 non-null   float64
7   density               1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates             1599 non-null   float64
10  alcohol               1599 non-null   float64
11  quality               1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

we can see all the columns are of the numerical data type

```
[5]: df.isnull().sum()
```

```
[5]: fixed_acidity          0
     volatile_acidity      0
     citric_acid           0
     residual_sugar        0
     chlorides             0
     free_sulfur_dioxide    0
     total_sulfur_dioxide   0
     density               0
     pH                   0
```

```

sulphates          0
alcohol            0
quality            0
dtype: int64

```

we are here checking for the null values in the dataset and we can say that there are no null values in the data

```

[6]: df1=df.copy()
     df1.head()

```

```

[6]:   fixed_acidity  volatile_acidity  citric_acid  residual_sugar  chlorides \
0           7.4           0.70           0.00           1.9           0.076
1           7.8           0.88           0.00           2.6           0.098
2           7.8           0.76           0.04           2.3           0.092
3          11.2           0.28           0.56           1.9           0.075
4           7.4           0.70           0.00           1.9           0.076

      free_sulfur_dioxide  total_sulfur_dioxide  density  pH  sulphates \
0              11.0              34.0  0.9978  3.51           0.56
1              25.0              67.0  0.9968  3.20           0.68
2              15.0              54.0  0.9970  3.26           0.65
3              17.0              60.0  0.9980  3.16           0.58
4              11.0              34.0  0.9978  3.51           0.56

      alcohol  quality
0         9.4         5
1         9.8         5
2         9.8         5
3         9.8         6
4         9.4         5

```

making a copy of the original data such that the operations done does not affect the original data

```

[7]: df1.drop('quality', inplace=True, axis=1)

```

here, we are dropping the target column quality

```

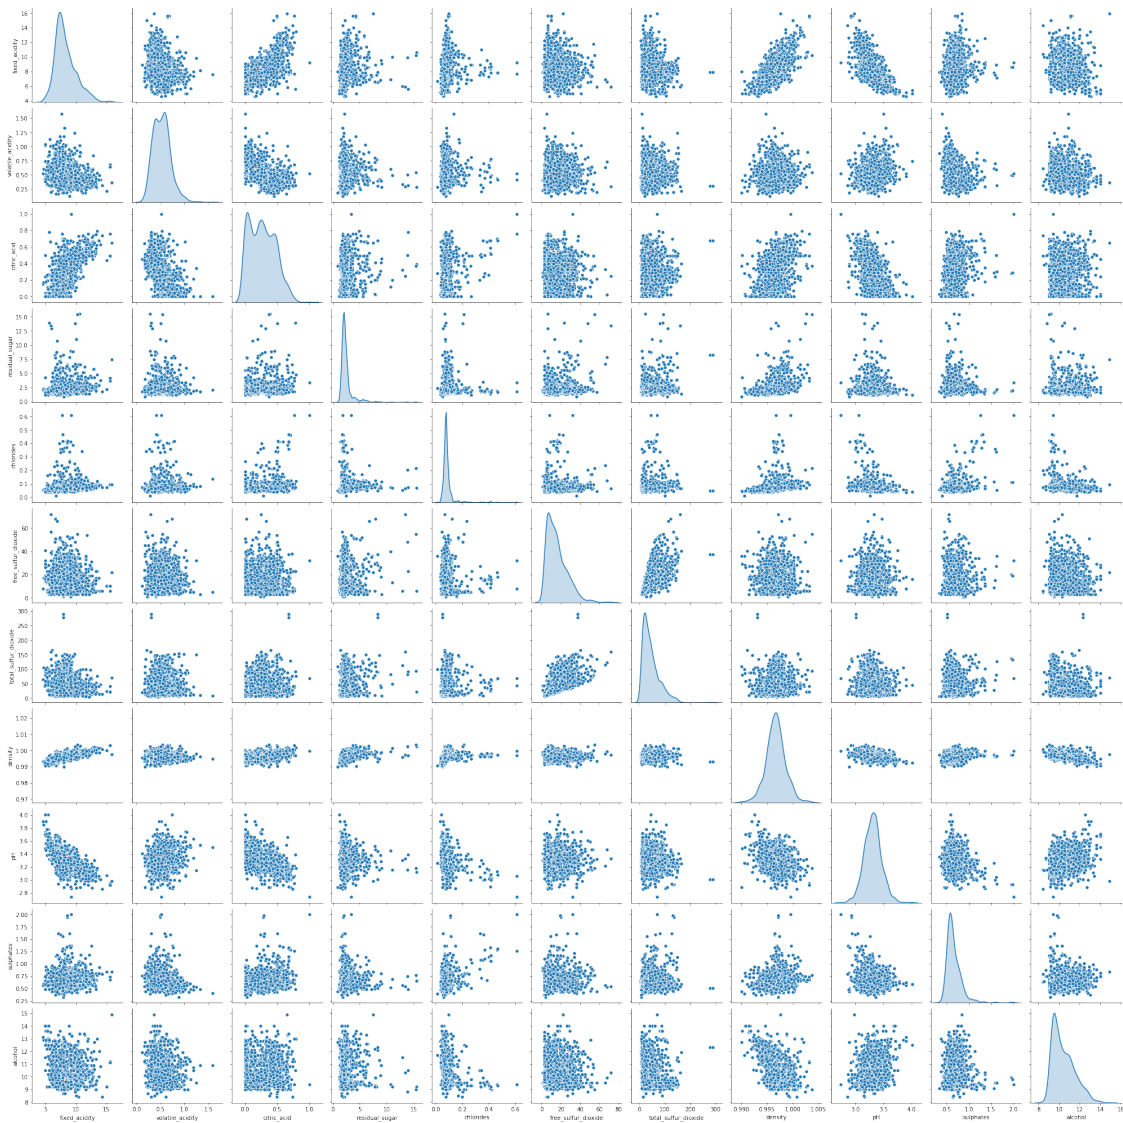
[8]: import seaborn as sns
     sns.pairplot(df1,diag_kind='kde')

```

```

[8]: <seaborn.axisgrid.PairGrid at 0x1638f678f60>

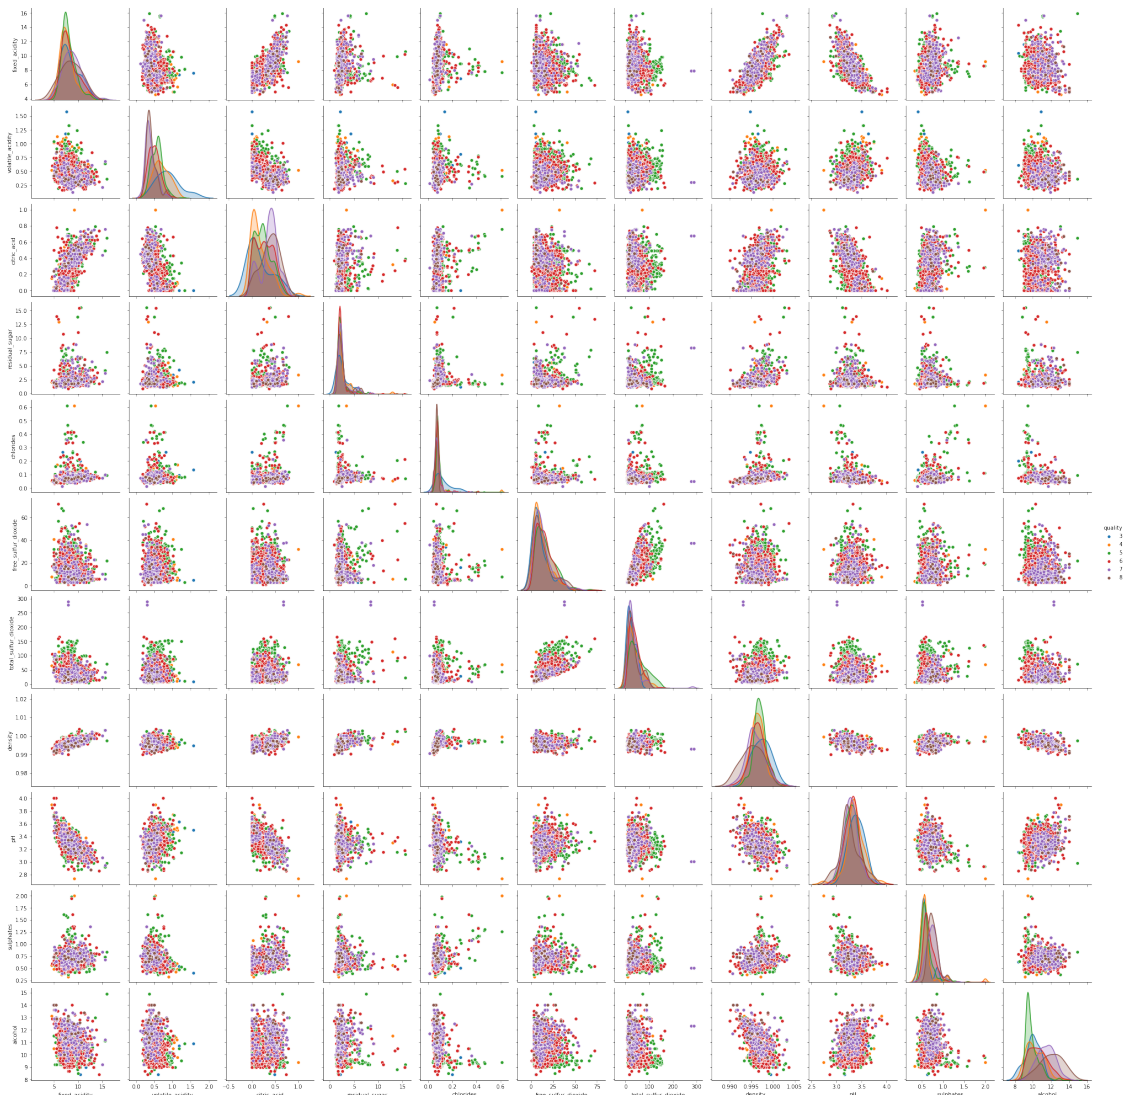
```



by seeing the pairplot we can see the kernel density plots and here if we see the small bumps on the curves which shows number of clusters

```
[41]: sns.pairplot(df,diag_kind='kde', hue='quality')
```

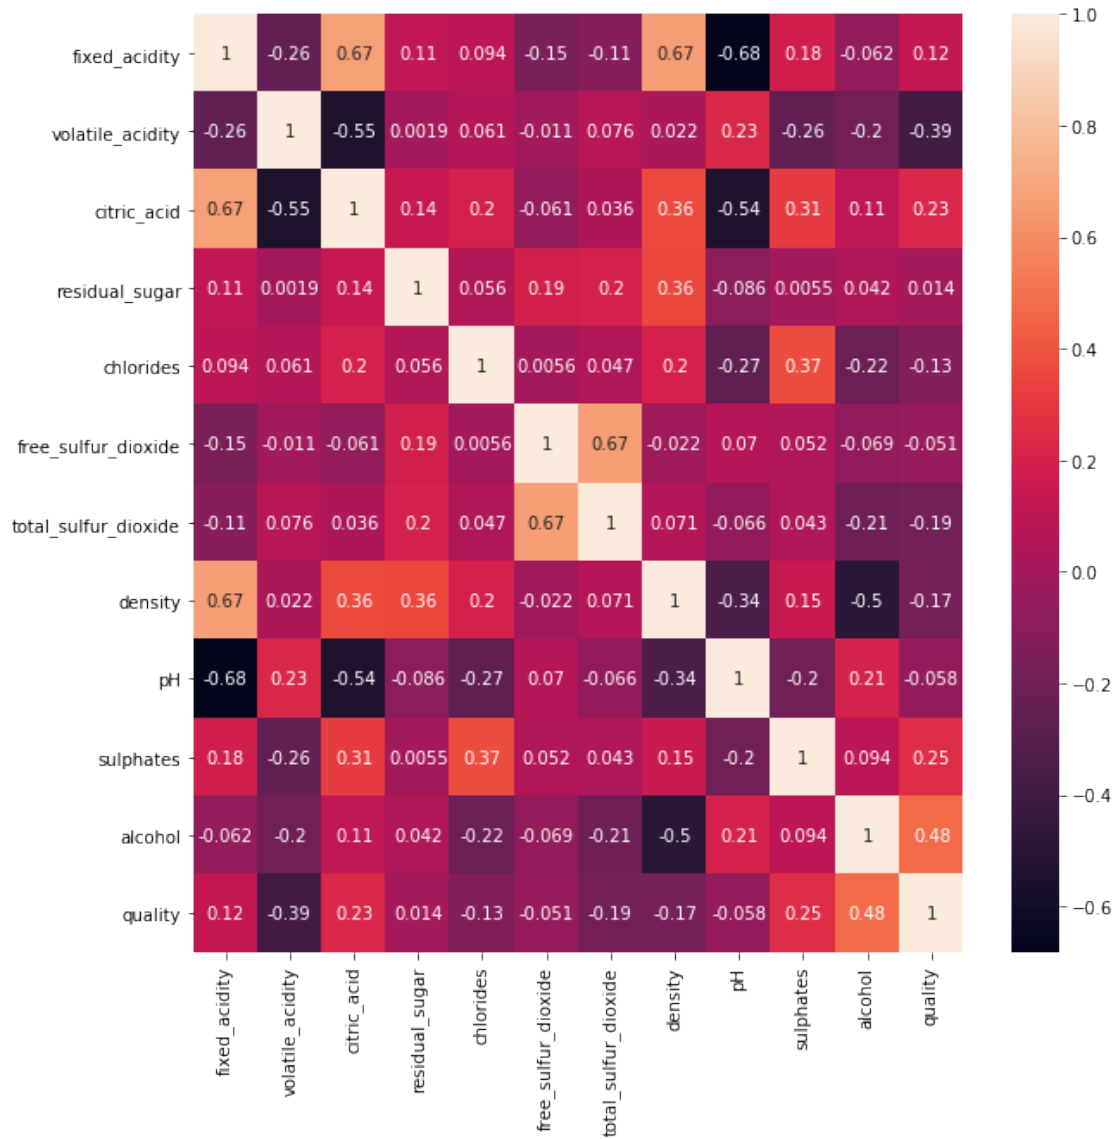
```
[41]: <seaborn.axisgrid.PairGrid at 0x16396b17550>
```



the pairplot includes the target variable where we can see the bumps on the curves as the number of clusters.

```
[15]: plt.figure(figsize=(10,10))
      sns.heatmap(df1.corr(), annot=True)
```

```
[15]: <matplotlib.axes._subplots.AxesSubplot at 0x26804021a90>
```



this is a corelation heat map of all the columns fixed acidity,citric acid and density, fixed acidity are having the highest positive correlation, citric acid, ph and fixed acidity,ph are having the strong negative correlation.

```
[15]: df1_scaled = df1.apply(zscore)
      df1_scaled.head()
```

```
[15]:   fixed_acidity  volatile_acidity  citric_acid  residual_sugar  chlorides  \
0      -0.528360           0.961877   -1.391472    -0.453218   -0.243707
1      -0.298547           1.967442   -1.391472     0.043416    0.223875
2      -0.298547           1.297065   -1.186070    -0.169427    0.096353
3       1.654856          -1.384443    1.484154    -0.453218   -0.264960
```

```

4      -0.528360          0.961877      -1.391472          -0.453218  -0.243707

      free_sulfur_dioxide  total_sulfur_dioxide  density      pH  sulphates  \
0      -0.466193          -0.379133  0.558274  1.288643  -0.579207
1       0.872638           0.624363  0.028261 -0.719933   0.128950
2     -0.083669           0.229047  0.134264 -0.331177  -0.048089
3       0.107592           0.411500  0.664277 -0.979104  -0.461180
4     -0.466193          -0.379133  0.558274  1.288643  -0.579207

      alcohol
0 -0.960246
1 -0.584777
2 -0.584777
3 -0.584777
4 -0.960246

```

we are here standardizing the data using the z score

```
[16]: model = KMeans(n_clusters = 3)
      model
```

```
[16]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
             n_clusters=3, n_init=10, n_jobs=None, precompute_distances='auto',
             random_state=None, tol=0.0001, verbose=0)
```

we are building a simple model by using the number of clusters as 3

```
[25]: cluster_range = range( 1, 15 )
      cluster_errors = []
      for num_clusters in cluster_range:
          clusters = KMeans( num_clusters, n_init = 10 )
          clusters.fit(df1_scaled)
          labels = clusters.labels_
          centroids = clusters.cluster_centers_
          cluster_errors.append( clusters.inertia_ )
      clusters_df = pd.DataFrame( { "num_clusters":cluster_range, "cluster_errors":_
      ↪cluster_errors } )
      clusters_df[0:15]
```

```
[25]:   num_clusters  cluster_errors
0         1      17589.000000
1         2      14330.119811
2         3      12629.906396
3         4      11294.409117
4         5      10155.374026
5         6       9362.868010
6         7       8644.994155
7         8       8299.254985
```

8	9	7966.196551
9	10	7699.666440
10	11	7451.519333
11	12	7215.884721
12	13	7064.082548
13	14	6815.347782

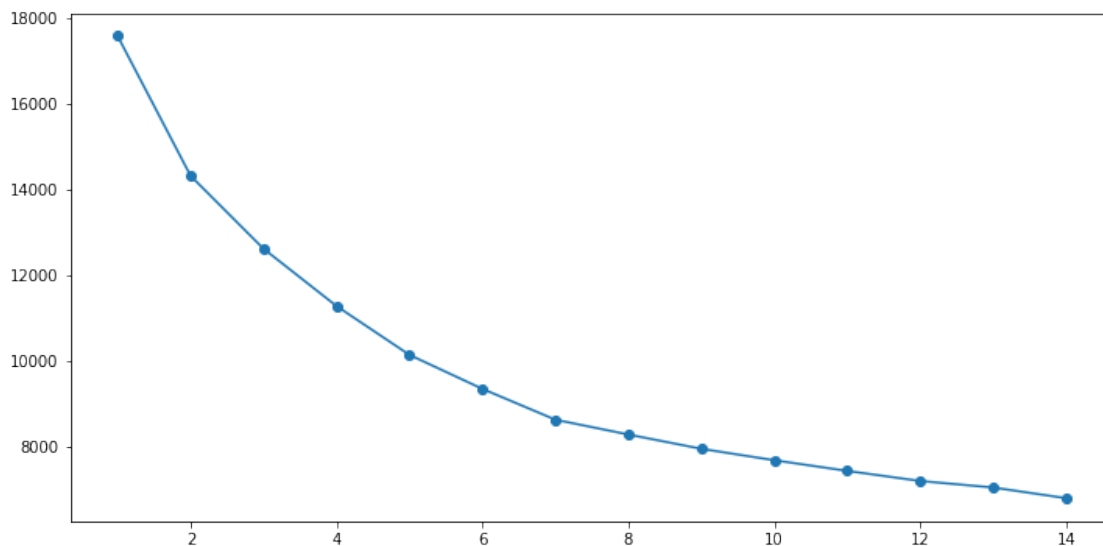
we are building model based on 3 clusters and also we are calculating the cluster errors

The total sum of squared distances of every data point from respective centroid is also called inertia. Let us print the inertia value for all K values. That K at which the inertia stop to drop significantly (elbow method) will be the best K.

```
[26]: # Elbow plot

plt.figure(figsize=(12,6))
plt.plot( clusters_df.num_clusters, clusters_df.cluster_errors, marker = "o" )
```

```
[26]: [<matplotlib.lines.Line2D at 0x16396b56e10>]
```



the elbow plot gives us the number of clusters to be used as we can see that the curve is slightly bent at 6 so we consider the number of clusters as 6 to build our model

```
[27]: kmeans = KMeans(n_clusters=6, n_init = 15, random_state=2345)
```

```
[28]: kmeans.fit(df1_scaled)
```

```
[28]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
            n_clusters=6, n_init=15, n_jobs=None, precompute_distances='auto',
            random_state=2345, tol=0.0001, verbose=0)
```



```
[29]: centroids = kmeans.cluster_centers_
centroids
```

```
[29]: array([[ -6.72451192e-02,  5.62073930e-02,  6.60225383e-02,
-1.00502756e-02, -2.95157193e-02,  9.77087975e-01,
 1.20800508e+00,  2.48331007e-01, -1.26341795e-01,
-1.70026513e-01, -5.62505055e-01],
 [ 1.33586541e+00, -6.74700057e-01,  1.12695753e+00,
 7.46994437e-02, -7.46339267e-03, -5.67340445e-01,
-5.43792234e-01,  7.63188823e-01, -8.44712801e-01,
 3.49276696e-01,  1.77163635e-01],
 [-4.83625180e-01,  6.85964431e-01, -8.31615045e-01,
-1.97792812e-01, -6.32528528e-02, -4.35623522e-01,
-4.27241532e-01, -6.53210396e-02,  4.19133687e-01,
-4.02054671e-01, -4.27167312e-01],
 [ 9.54162998e-02,  2.19980305e-03,  1.18155266e+00,
-3.89872163e-01,  5.78475973e+00, -4.95156003e-02,
 5.10329601e-01,  1.80071833e-01, -1.73579154e+00,
 3.66341219e+00, -8.69731260e-01],
 [-6.92918829e-01, -4.37887177e-01, -1.47498832e-01,
-2.57759586e-01, -4.17044161e-01,  1.24533302e-01,
-2.29504460e-01, -1.24647482e+00,  6.33635100e-01,
 1.35008034e-01,  1.29557245e+00],
 [-8.56332101e-02, -3.46521643e-02,  4.14855742e-01,
 4.96176756e+00,  2.96387895e-01,  1.75019116e+00,
 1.69583213e+00,  1.22500051e+00, -3.25459600e-01,
-2.37893278e-02, -3.63912996e-01]])
```

```
[30]: centroid_df = pd.DataFrame(centroids, columns = list(df1_scaled) )
centroid_df
```

```
[30]:
```

	fixed_acidity	volatile_acidity	citric_acid	residual_sugar	chlorides	\
0	-0.067245	0.056207	0.066023	-0.010050	-0.029516	
1	1.335865	-0.674700	1.126958	0.074699	-0.007463	
2	-0.483625	0.685964	-0.831615	-0.197793	-0.063253	
3	0.095416	0.002200	1.181553	-0.389872	5.784760	
4	-0.692919	-0.437887	-0.147499	-0.257760	-0.417044	
5	-0.085633	-0.034652	0.414856	4.961768	0.296388	

	free_sulfur_dioxide	total_sulfur_dioxide	density	pH	sulphates	\
0	0.977088	1.208005	0.248331	-0.126342	-0.170027	
1	-0.567340	-0.543792	0.763189	-0.844713	0.349277	
2	-0.435624	-0.427242	-0.065321	0.419134	-0.402055	
3	-0.049516	0.510330	0.180072	-1.735792	3.663412	
4	0.124533	-0.229504	-1.246475	0.633635	0.135008	
5	1.750191	1.695832	1.225001	-0.325460	-0.023789	

```

    alcohol
0 -0.562505
1  0.177164
2 -0.427167
3 -0.869731
4  1.295572
5 -0.363913

```

```

[33]: ## creating a new dataframe only for labels and converting it into categorical
      ↪variable
df1_labels = pd.DataFrame(kmeans.labels_ , columns = list(['labels']))

df1_labels['labels'] = df1_labels['labels'].astype('category')

```

```

[34]: # Joining the label dataframe with the Wine data frame to create
      ↪wine_df_labeled. Note: it could be appended to original dataframe
snail_df1_labeled = df1.join(df1_labels)

```

```

[36]: df1_analysis = (snail_df1_labeled.groupby(['labels'] , axis=0)).head(4177) #
      ↪the groupby creates a groupeddataframe that needs
      # to be converted back to dataframe. I am using .head(30000) for that
df1_analysis

```

```

[36]:      fixed_acidity  volatile_acidity  citric_acid  residual_sugar  chlorides  \
0              7.4             0.700           0.00             1.9         0.076
1              7.8             0.880           0.00             2.6         0.098
2              7.8             0.760           0.04             2.3         0.092
3             11.2             0.280           0.56             1.9         0.075
4              7.4             0.700           0.00             1.9         0.076
...           ...                   ...           ...           ...           ...
1594            6.2             0.600           0.08             2.0         0.090
1595            5.9             0.550           0.10             2.2         0.062
1596            6.3             0.510           0.13             2.3         0.076
1597            5.9             0.645           0.12             2.0         0.075
1598            6.0             0.310           0.47             3.6         0.067

```

```

      free_sulfur_dioxide  total_sulfur_dioxide  density    pH  sulphates  \
0              11.0             34.0  0.99780  3.51         0.56
1              25.0             67.0  0.99680  3.20         0.68
2              15.0             54.0  0.99700  3.26         0.65
3              17.0             60.0  0.99800  3.16         0.58
4              11.0             34.0  0.99780  3.51         0.56
...           ...                   ...           ...           ...
1594            32.0             44.0  0.99490  3.45         0.58
1595            39.0             51.0  0.99512  3.52         0.76
1596            29.0             40.0  0.99574  3.42         0.75
1597            32.0             44.0  0.99547  3.57         0.71

```

1598	18.0	42.0	0.99549	3.39	0.66
------	------	------	---------	------	------

	alcohol	labels
0	9.4	2
1	9.8	2
2	9.8	2
3	9.8	1
4	9.4	2
...
1594	10.5	4
1595	11.2	4
1596	11.0	4
1597	10.2	2
1598	11.0	4

[1599 rows x 12 columns]

```
[37]: snail_df1_labeled['labels'].value_counts()
```

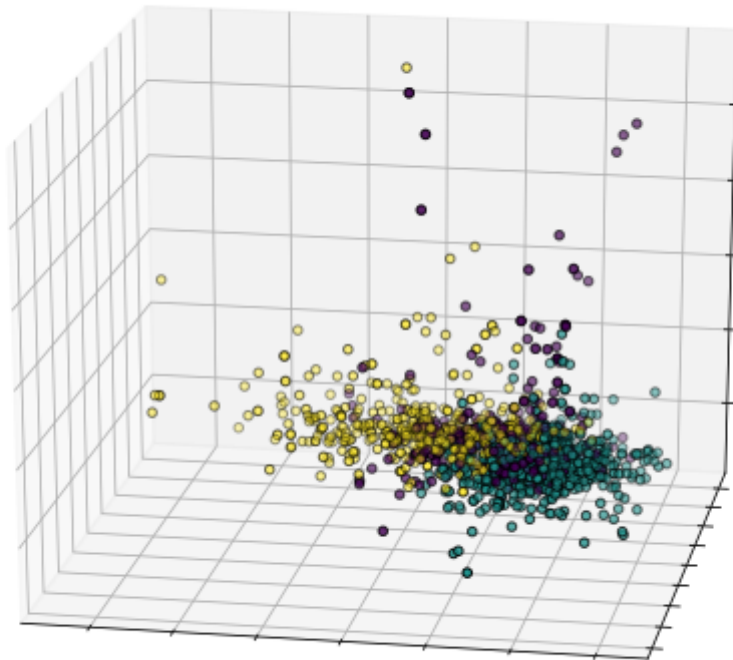
```
[37]: 2    524
      1    364
      0    347
      4    302
      5     34
      3     28
      Name: labels, dtype: int64
```

```
[38]: from mpl_toolkits.mplot3d import Axes3D
```

```
[61]: fig = plt.figure(figsize=(8, 6))
      ax = Axes3D(fig, rect=[0, 0, .95, 1], elev=20, azimuth=100)
      kmeans.fit(df1_scaled)
      labels = kmeans.labels_
      ax.scatter(df1_scaled.iloc[:, 0], df1_scaled.iloc[:, 1], df1_scaled.iloc[:, 2],
                  c=labels.astype(np.float), edgecolor='k')
      ax.w_xaxis.set_ticklabels([])
      ax.w_yaxis.set_ticklabels([])
      ax.w_zaxis.set_ticklabels([])
      ax.set_title('3D plot of KMeans Clustering')
```

```
[61]: Text(0.5, 0.92, '3D plot of KMeans Clustering')
```

3D plot of KMeans Clustering



Use features fixed acidity and volatile acidity

```
[44]: df2 = df1_scaled.loc[:, 'fixed_acidity': 'volatile_acidity']  
df2.head()
```

```
[44]:   fixed_acidity  volatile_acidity  
0      -0.528360         0.961877  
1      -0.298547         1.967442  
2      -0.298547         1.297065  
3       1.654856        -1.384443  
4      -0.528360         0.961877
```

```
[45]: model = KMeans(n_clusters = 3)  
model
```

```
[45]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,  
         n_clusters=3, n_init=10, n_jobs=None, precompute_distances='auto',  
         random_state=None, tol=0.0001, verbose=0)
```

```
[49]: cluster_range = range( 1, 15 )
cluster_errors = []
for num_clusters in cluster_range:
    clusters = KMeans( num_clusters, n_init = 10 )
    clusters.fit(df2)
    labels = clusters.labels_
    centroids = clusters.cluster_centers_
    cluster_errors.append( clusters.inertia_ )
clusters_df = pd.DataFrame( { "num_clusters":cluster_range, "cluster_errors":_
↪cluster_errors } )
clusters_df[0:15]
```

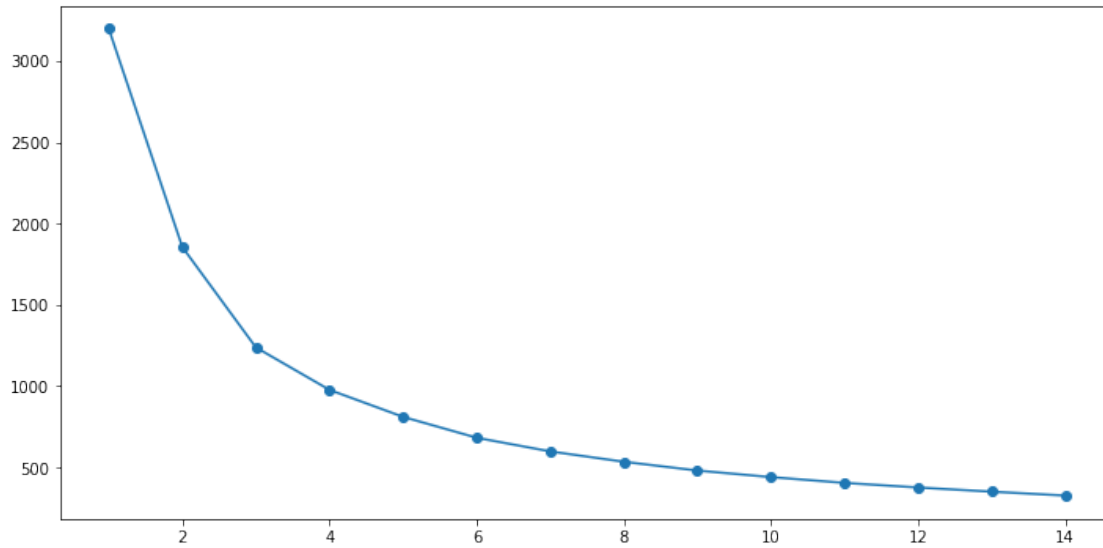
```
[49]:
```

	num_clusters	cluster_errors
0	1	3198.000000
1	2	1855.089000
2	3	1237.886499
3	4	976.934633
4	5	810.980586
5	6	682.546130
6	7	598.547803
7	8	534.821305
8	9	480.822240
9	10	440.501462
10	11	404.879042
11	12	376.540714
12	13	350.924509
13	14	326.937199

```
[50]: # Elbow plot

plt.figure(figsize=(12,6))
plt.plot( clusters_df.num_clusters, clusters_df.cluster_errors, marker = "o" )
```

```
[50]: [<matplotlib.lines.Line2D at 0x1639cc142e8>]
```



the value of k is 3 in this case

```
[51]: kmeans = KMeans(n_clusters=3, n_init = 15, random_state=2345)
```

```
[52]: kmeans.fit(df2)
```

```
[52]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
            n_clusters=3, n_init=15, n_jobs=None, precompute_distances='auto',
            random_state=2345, tol=0.0001, verbose=0)
```

```
[53]: centroids = kmeans.cluster_centers_
centroids
```

```
[53]: array([[ -0.43343719,  0.93958281],
            [ 1.47529531, -0.60241457],
            [-0.42029776, -0.6523901 ]])
```

```
[55]: centroid_df = pd.DataFrame(centroids, columns = list(df2) )
centroid_df
```

```
[55]:    fixed_acidity  volatile_acidity
0      -0.433437      0.939583
1       1.475295     -0.602415
2      -0.420298     -0.652390
```

```
[56]: ## creating a new dataframe only for labels and converting it into categorical
      ↪ variable
df_labels = pd.DataFrame(kmeans.labels_ , columns = list(['labels']))
```

```
df_labels['labels'] = df_labels['labels'].astype('category')
```

```
[57]: snail_df_labeled = df2.join(df_labels)
```

```
[58]: df_analysis = (snail_df_labeled.groupby(['labels'] , axis=0)).head(4177) # the
↳ groupby creates a groupeddataframe that needs
# to be converted back to dataframe. I am using .head(30000) for that
df_analysis
```

```
[58]:
```

	fixed_acidity	volatile_acidity	labels
0	-0.528360	0.961877	0
1	-0.298547	1.967442	0
2	-0.298547	1.297065	0
3	1.654856	-1.384443	1
4	-0.528360	0.961877	0
...
1594	-1.217796	0.403229	0
1595	-1.390155	0.123905	2
1596	-1.160343	-0.099554	2
1597	-1.390155	0.654620	0
1598	-1.332702	-1.216849	2

[1599 rows x 3 columns]

```
[59]: snail_df_labeled['labels'].value_counts()
```

```
[59]: 0    644
      2    596
      1    359
      Name: labels, dtype: int64
```

Use features Citric acidity and fixed acidity

```
[72]: df3 = df1_scaled[['citric_acid','fixed_acidity']]
      df3.head()
```

```
[72]:
```

	citric_acid	fixed_acidity
0	-1.391472	-0.528360
1	-1.391472	-0.298547
2	-1.186070	-0.298547
3	1.484154	1.654856
4	-1.391472	-0.528360

```
[73]: model = KMeans(n_clusters = 3)
      model
```

```
[73]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
            n_clusters=3, n_init=10, n_jobs=None, precompute_distances='auto',
            random_state=None, tol=0.0001, verbose=0)
```

```
[74]: cluster_range = range( 1, 15 )
cluster_errors = []
for num_clusters in cluster_range:
    clusters = KMeans( num_clusters, n_init = 10 )
    clusters.fit(df3)
    # labels = clusters.labels_
    # centroids = clusters.cluster_centers_
    cluster_errors.append( clusters.inertia_ )
clusters_df = pd.DataFrame( { "num_clusters":cluster_range, "cluster_errors":
    ↪cluster_errors } )
clusters_df[0:15]
```

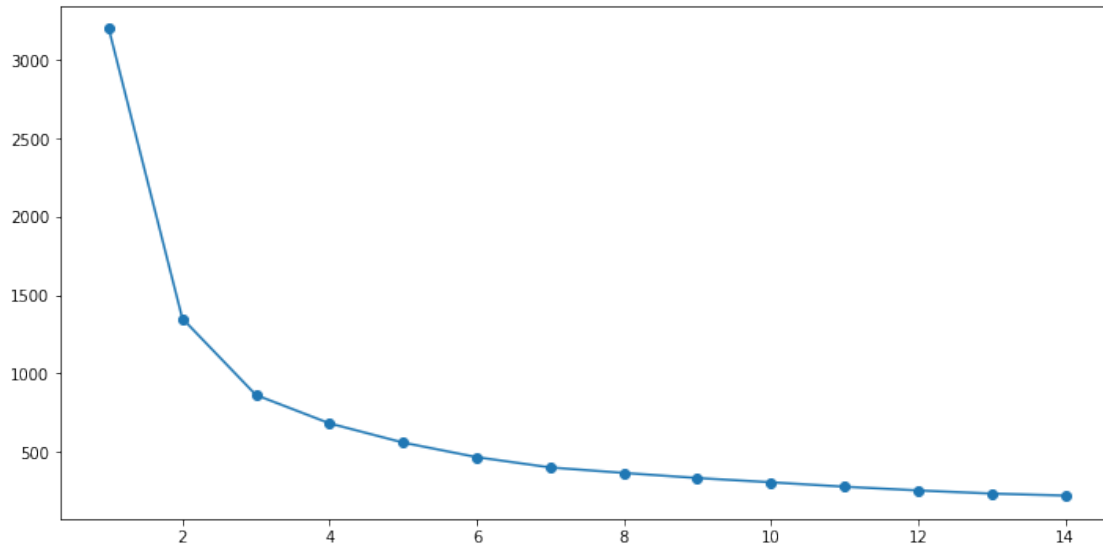
```
[74]:
```

	num_clusters	cluster_errors
0	1	3198.000000
1	2	1349.582779
2	3	861.704601
3	4	680.668631
4	5	558.073517
5	6	465.071896
6	7	399.431950
7	8	364.434654
8	9	331.892906
9	10	304.718267
10	11	276.531927
11	12	252.743949
12	13	232.644904
13	14	220.003736

```
[75]: # Elbow plot

plt.figure(figsize=(12,6))
plt.plot( clusters_df.num_clusters, clusters_df.cluster_errors, marker = "o" )
```

```
[75]: [<matplotlib.lines.Line2D at 0x1639ce3e438>]
```

1 here we are taking the value of $k=3$

```
[76]: kmeans = KMeans(n_clusters=3, n_init = 15, random_state=2345)
```

```
[77]: kmeans.fit(df3)
```

```
[77]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
            n_clusters=3, n_init=15, n_jobs=None, precompute_distances='auto',
            random_state=2345, tol=0.0001, verbose=0)
```

```
[79]: centroids = kmeans.cluster_centers_
centroids
```

```
[79]: array([[ -0.95846298, -0.70641229],
            [ 0.38702498, -0.0232819 ],
            [ 1.30809487,  1.61190945]])
```

```
[81]: centroid_df = pd.DataFrame(centroids, columns = list(df3) )
centroid_df
```

```
[81]:   citric_acid  fixed_acidity
0    -0.958463    -0.706412
1     0.387025    -0.023282
2     1.308095     1.611909
```

```
[82]: ## creating a new dataframe only for labels and converting it into categorical
      ↪ variable
```

```
df_labels = pd.DataFrame(kmeans.labels_ , columns = list(['labels']))

df_labels['labels'] = df_labels['labels'].astype('category')
```

```
[83]: snail_df_labeled = df3.join(df_labels)
```

```
[84]: df_analysis = (snail_df_labeled.groupby(['labels'] , axis=0)).head(4177) # the
      ↳ groupby creates a groupeddataframe that needs
      # to be converted back to dataframe. I am using .head(30000) for that
df_analysis
```

```
[84]:      citric_acid  fixed_acidity labels
0      -1.391472      -0.528360      0
1      -1.391472      -0.298547      0
2      -1.186070      -0.298547      0
3       1.484154       1.654856      2
4      -1.391472      -0.528360      0
...
1594    -0.980669     -1.217796      0
1595    -0.877968     -1.390155      0
1596    -0.723916     -1.160343      0
1597    -0.775267     -1.390155      0
1598     1.021999     -1.332702      1
```

[1599 rows x 3 columns]

```
[85]: snail_df_labeled['labels'].value_counts()
```

```
[85]: 0    666
      1    632
      2    301
      Name: labels, dtype: int64
```

```
[ ]: #3 Use feautes residual suger and sulphades
```

```
[88]: df.columns
```

```
[88]: Index(['fixed_acidity', 'volatile_acidity', 'citric_acid', 'residual_sugar',
          'chlorides', 'free_sulfur_dioxide', 'total_sulfur_dioxide', 'density',
          'pH', 'sulphates', 'alcohol', 'quality'],
          dtype='object')
```

```
[91]: df4 = df1_scaled[['residual_sugar', 'sulphates']]
      df4.head()
```

```
[91]:      residual_sugar  sulphates
0      -0.453218   -0.579207
```

```

1      0.043416   0.128950
2     -0.169427  -0.048089
3     -0.453218  -0.461180
4     -0.453218  -0.579207

```

```
[92]: model = KMeans(n_clusters = 3)
      model
```

```
[92]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
            n_clusters=3, n_init=10, n_jobs=None, precompute_distances='auto',
            random_state=None, tol=0.0001, verbose=0)
```

```
[93]: cluster_range = range( 1, 15 )
      cluster_errors = []
      for num_clusters in cluster_range:
          clusters = KMeans( num_clusters, n_init = 10 )
          clusters.fit(df4)
          # labels = clusters.labels_
          # centroids = clusters.cluster_centers_
          cluster_errors.append( clusters.inertia_ )
      clusters_df = pd.DataFrame( { "num_clusters":cluster_range, "cluster_errors":_
      ↪cluster_errors } )
      clusters_df[0:15]
```

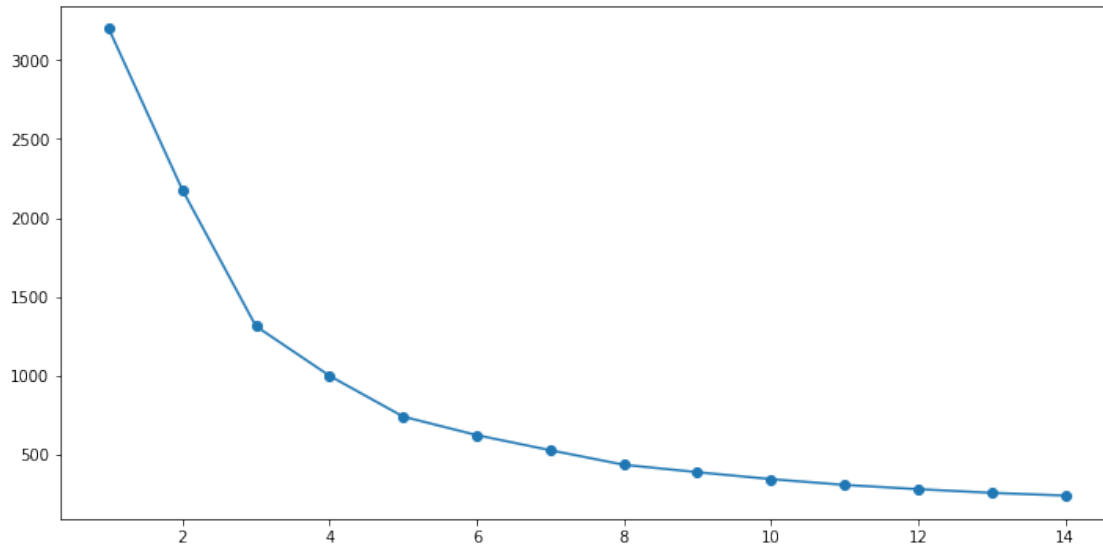
```
[93]:
```

	num_clusters	cluster_errors
0	1	3198.000000
1	2	2177.114997
2	3	1312.577756
3	4	998.173046
4	5	739.725164
5	6	622.227993
6	7	526.452113
7	8	434.267238
8	9	387.224194
9	10	343.239636
10	11	306.545590
11	12	279.599023
12	13	256.186147
13	14	238.970926

```
[94]: # Elbow plot

      plt.figure(figsize=(12,6))
      plt.plot( clusters_df.num_clusters, clusters_df.cluster_errors, marker = "o" )
```

```
[94]: [<matplotlib.lines.Line2D at 0x1639ceab4e0>]
```



here we are taking k value as 5

```
[97]: kmeans = KMeans(n_clusters=3, n_init = 15, random_state=2345)
```

```
[99]: kmeans.fit(df4)
```

```
[99]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
             n_clusters=3, n_init=15, n_jobs=None, precompute_distances='auto',
             random_state=2345, tol=0.0001, verbose=0)
```

```
[100]: centroids = kmeans.cluster_centers_
        centroids
```

```
[100]: array([[ -0.18614032, -0.40542353],
              [-0.19888324,  1.41705131],
              [ 3.36585956,  0.01717269]])
```

```
[102]: centroid_df = pd.DataFrame(centroids, columns = list(df4) )
        centroid_df
```

```
[102]:   residual_sugar  sulphates
0      -0.186140  -0.405424
1      -0.198883   1.417051
2       3.365860   0.017173
```

```
[106]: ## creating a new dataframe only for labels and converting it into categorical
        ↪ variable
        df_labels = pd.DataFrame(kmeans.labels_ , columns = list(['labels']))
```

```
df_labels['labels'] = df_labels['labels'].astype('category')
```

```
[107]: # Joining the label dataframe with the Wine data frame to create
       ↪ wine_df_labeled. Note: it could be appended to original dataframe
       snail_df_labeled = df2.join(df_labels)
```

```
[108]: df_analysis = (snail_df_labeled.groupby(['labels'] , axis=0)).head(4177) # the
       ↪ groupby creates a groupeddataframe that needs
       # to be converted back to dataframe. I am using .head(30000) for that
       df_analysis
```

```
[108]:      labels
0         0
1         0
2         0
3         0
4         0
...      ...
1594      0
1595      1
1596      1
1597      0
1598      0

[1599 rows x 1 columns]
```

```
[109]: snail_df_labeled['labels'].value_counts()
```

```
[109]: 0    1178
       1     336
       2      85
       Name: labels, dtype: int64
```

#4 Use features free.sulfur.dioxide and total.sulfur.dioxide

```
[88]: df.columns
```

```
[88]: Index(['fixed_acidity', 'volatile_acidity', 'citric_acid', 'residual_sugar',
         'chlorides', 'free_sulfur_dioxide', 'total_sulfur_dioxide', 'density',
         'pH', 'sulphates', 'alcohol', 'quality'],
         dtype='object')
```

```
[112]: df5 = df1_scaled[['free_sulfur_dioxide', 'total_sulfur_dioxide']]
       df5.head()
```

```
[112]:   free_sulfur_dioxide  total_sulfur_dioxide
0             -0.466193             -0.379133
```

1	0.872638	0.624363
2	-0.083669	0.229047
3	0.107592	0.411500
4	-0.466193	-0.379133

```
[113]: model = KMeans(n_clusters = 3)
model
```

```
[113]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
             n_clusters=3, n_init=10, n_jobs=None, precompute_distances='auto',
             random_state=None, tol=0.0001, verbose=0)
```

```
[114]: cluster_range = range( 1, 15 )
cluster_errors = []
for num_clusters in cluster_range:
    clusters = KMeans( num_clusters, n_init = 10 )
    clusters.fit(df5)
    # labels = clusters.labels_
    # centroids = clusters.cluster_centers_
    cluster_errors.append( clusters.inertia_ )
clusters_df = pd.DataFrame( { "num_clusters":cluster_range, "cluster_errors":_
    ↪cluster_errors } )
clusters_df[0:15]
```

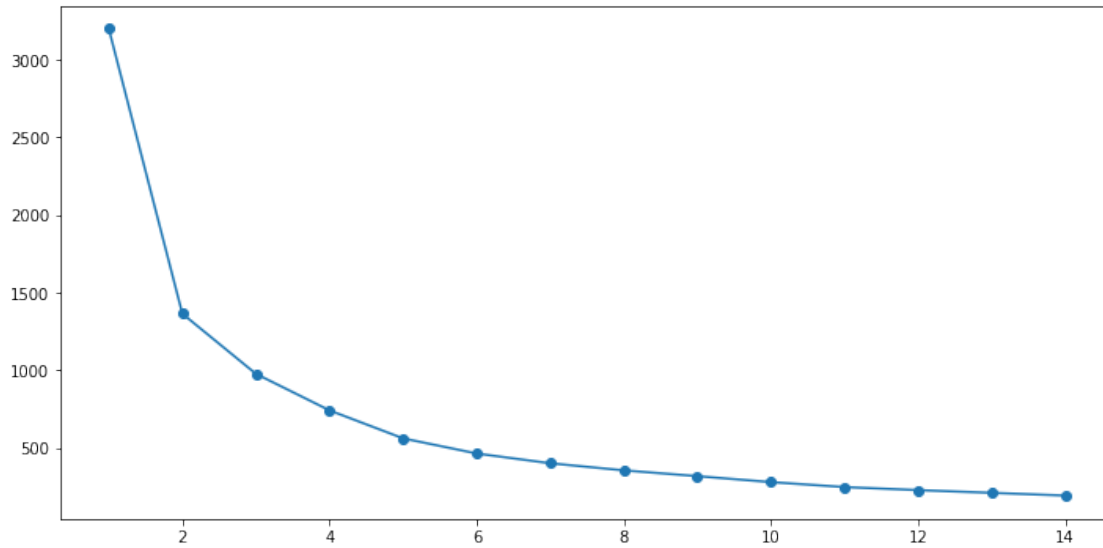
```
[114]:
```

	num_clusters	cluster_errors
0	1	3198.000000
1	2	1363.149246
2	3	975.837335
3	4	740.680585
4	5	560.417140
5	6	462.895676
6	7	400.038118
7	8	354.388299
8	9	317.066485
9	10	278.740024
10	11	246.775241
11	12	227.159562
12	13	209.648403
13	14	192.081726

```
[115]: # Elbow plot

plt.figure(figsize=(12,6))
plt.plot( clusters_df.num_clusters, clusters_df.cluster_errors, marker = "o" )
```

```
[115]: [<matplotlib.lines.Line2D at 0x1639ced5e10>]
```



here we are taking k value as 2

```
[116]: kmeans = KMeans(n_clusters=3, n_init = 15, random_state=2345)
```

```
[117]: kmeans.fit(df4)
```

```
[117]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
              n_clusters=3, n_init=15, n_jobs=None, precompute_distances='auto',
              random_state=2345, tol=0.0001, verbose=0)
```

```
[118]: centroids = kmeans.cluster_centers_
centroids
```

```
[118]: array([[ -0.18614032, -0.40542353],
              [-0.19888324,  1.41705131],
              [ 3.36585956,  0.01717269]])
```

```
[119]: centroid_df = pd.DataFrame(centroids, columns = list(df4) )
centroid_df
```

```
[119]:   residual_sugar  sulphates
0      -0.186140  -0.405424
1      -0.198883   1.417051
2       3.365860   0.017173
```

```
[120]: ## creating a new dataframe only for labels and converting it into categorical
      ↪ variable
df_labels = pd.DataFrame(kmeans.labels_ , columns = list(['labels']))
```

```
df_labels['labels'] = df_labels['labels'].astype('category')
```

```
[121]: # Joining the label dataframe with the Wine data frame to create
       ↪ wine_df_labeled. Note: it could be appended to original dataframe
       snail_df_labeled = df2.join(df_labels)
```

```
[122]: df_analysis = (snail_df_labeled.groupby(['labels'] , axis=0)).head(4177) # the
       ↪ groupby creates a groupeddataframe that needs
       # to be converted back to dataframe. I am using .head(30000) for that
       df_analysis
```

```
[122]:      labels
0         0
1         0
2         0
3         0
4         0
...      ...
1594      0
1595      1
1596      1
1597      0
1598      0

[1599 rows x 1 columns]
```

```
[123]: snail_df_labeled['labels'].value_counts()
```

```
[123]: 0    1178
       1     336
       2      85
       Name: labels, dtype: int64
```

Use features fixed acidity, citric acidity and volatile acidity

```
[ ]:
```

```
[88]: df.columns
```

```
[88]: Index(['fixed_acidity', 'volatile_acidity', 'citric_acid', 'residual_sugar',
        'chlorides', 'free_sulfur_dioxide', 'total_sulfur_dioxide', 'density',
        'pH', 'sulphates', 'alcohol', 'quality'],
        dtype='object')
```

```
[127]: df6 = df1_scaled[['fixed_acidity', 'volatile_acidity', 'citric_acid']]
       df6.head()
```



```
[127]:    fixed_acidity  volatile_acidity  citric_acid
0      -0.528360          0.961877    -1.391472
1      -0.298547          1.967442    -1.391472
2      -0.298547          1.297065    -1.186070
3       1.654856         -1.384443     1.484154
4      -0.528360          0.961877    -1.391472
```

```
[128]: model = KMeans(n_clusters = 3)
model
```

```
[128]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
             n_clusters=3, n_init=10, n_jobs=None, precompute_distances='auto',
             random_state=None, tol=0.0001, verbose=0)
```

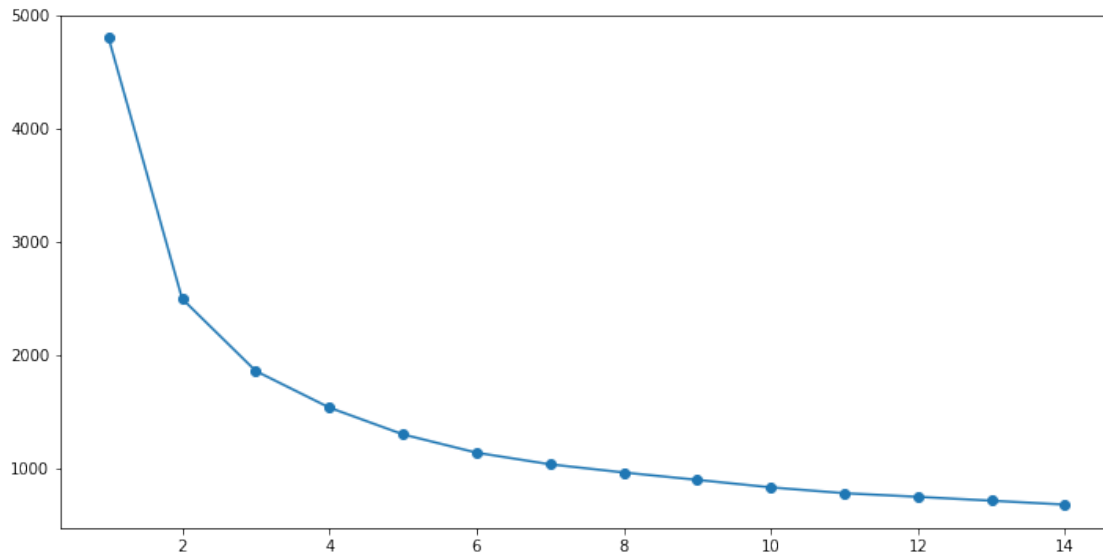
```
[129]: cluster_range = range( 1, 15 )
cluster_errors = []
for num_clusters in cluster_range:
    clusters = KMeans( num_clusters, n_init = 10 )
    clusters.fit(df6)
    # labels = clusters.labels_
    # centroids = clusters.cluster_centers_
    cluster_errors.append( clusters.inertia_ )
clusters_df = pd.DataFrame( { "num_clusters":cluster_range, "cluster_errors":_
    ↪cluster_errors } )
clusters_df[0:15]
```

```
[129]:    num_clusters  cluster_errors
0         1      4797.000000
1         2      2494.155753
2         3      1857.127760
3         4      1534.077522
4         5      1297.582734
5         6      1135.794701
6         7      1033.152006
7         8       961.045212
8         9       895.923730
9        10       829.524099
10       11       777.970401
11       12       746.043602
12       13       711.805290
13       14       677.369459
```

```
[130]: # Elbow plot

plt.figure(figsize=(12,6))
plt.plot( clusters_df.num_clusters, clusters_df.cluster_errors, marker = "o" )
```

```
[130]: [<matplotlib.lines.Line2D at 0x1639cf46390>]
```



here we are taking k value as 2

```
[131]: kmeans = KMeans(n_clusters=3, n_init = 15, random_state=2345)
```

```
[133]: kmeans.fit(df6)
```

```
[133]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,  
          n_clusters=3, n_init=15, n_jobs=None, precompute_distances='auto',  
          random_state=2345, tol=0.0001, verbose=0)
```

```
[134]: centroids = kmeans.cluster_centers_  
centroids
```

```
[134]: array([[ -0.51869461,  0.76068147, -0.83114207],  
        [ 1.5514048 , -0.60938513,  1.27619223],  
        [-0.20699183, -0.70725481,  0.40239642]])
```

```
[136]: centroid_df = pd.DataFrame(centroids, columns = list(df6) )  
centroid_df
```

```
[136]:    fixed_acidity  volatile_acidity  citric_acid  
0      -0.518695         0.760681    -0.831142  
1       1.551405        -0.609385     1.276192  
2      -0.206992        -0.707255     0.402396
```

```
[137]: ## creating a new dataframe only for labels and converting it into categorical
      ↪variable
      df_labels = pd.DataFrame(kmeans.labels_ , columns = list(['labels']))

      df_labels['labels'] = df_labels['labels'].astype('category')
```

```
[138]: # Joining the label dataframe with the Wine data frame to create
      ↪wine_df_labeled. Note: it could be appended to original dataframe
      snail_df_labeled = df2.join(df_labels)
```

```
[139]: df_analysis = (snail_df_labeled.groupby(['labels'] , axis=0)).head(4177) # the
      ↪groupby creates a groupeddataframe that needs
      # to be converted back to dataframe. I am using .head(30000) for that
      df_analysis
```

```
[139]:      labels
0         0
1         0
2         0
3         1
4         0
...
1594      0
1595      0
1596      0
1597      0
1598      2

[1599 rows x 1 columns]
```

```
[140]: snail_df_labeled['labels'].value_counts()
```

```
[140]: 0    749
      2    528
      1    322
      Name: labels, dtype: int64
```

```
[ ]:
```

```
[ ]:
```

#4 Use features density, ph

```
[88]: df.columns
```

```
[88]: Index(['fixed_acidity', 'volatile_acidity', 'citric_acid', 'residual_sugar',
      'chlorides', 'free_sulfur_dioxide', 'total_sulfur_dioxide', 'density',
```

```
    'pH', 'sulphates', 'alcohol', 'quality'],  
    dtype='object')
```

```
[141]: df7 = df1_scaled[[ 'density', 'pH']]  
df7.head()
```

```
[141]:      density      pH  
0  0.558274  1.288643  
1  0.028261 -0.719933  
2  0.134264 -0.331177  
3  0.664277 -0.979104  
4  0.558274  1.288643
```

```
[142]: model = KMeans(n_clusters = 3)  
model
```

```
[142]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,  
            n_clusters=3, n_init=10, n_jobs=None, precompute_distances='auto',  
            random_state=None, tol=0.0001, verbose=0)
```

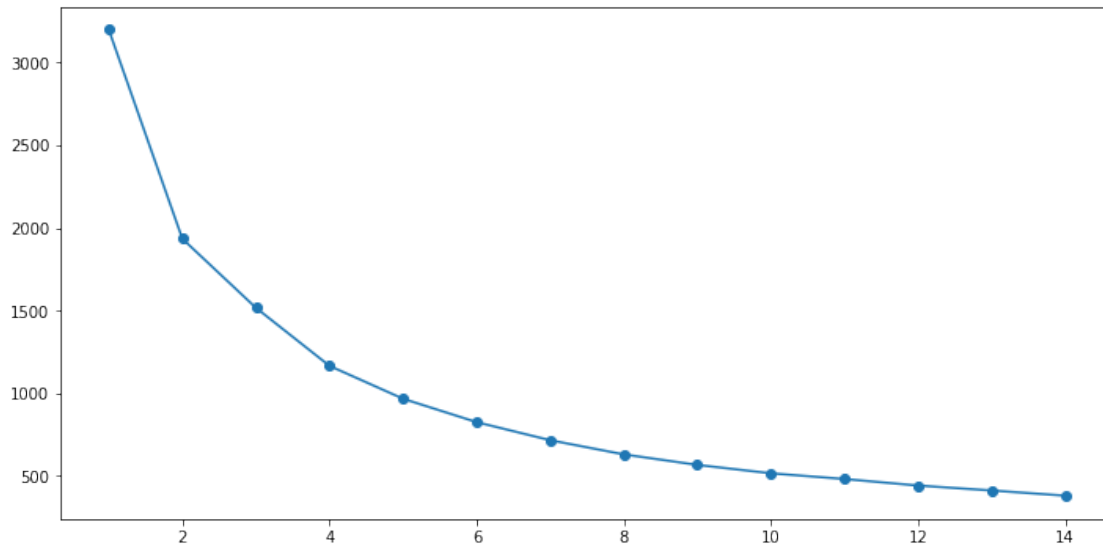
```
[143]: cluster_range = range( 1, 15 )  
cluster_errors = []  
for num_clusters in cluster_range:  
    clusters = KMeans( num_clusters, n_init = 10 )  
    clusters.fit(df7)  
    # labels = clusters.labels_  
    # centroids = clusters.cluster_centers_  
    cluster_errors.append( clusters.inertia_ )  
clusters_df = pd.DataFrame( { "num_clusters":cluster_range, "cluster_errors":  
    ↪cluster_errors } )  
clusters_df[0:15]
```

```
[143]:      num_clusters  cluster_errors  
0           1      3198.000000  
1           2      1934.185304  
2           3      1516.801912  
3           4      1165.674841  
4           5       966.910641  
5           6       825.483154  
6           7       716.252814  
7           8       630.578389  
8           9       567.017297  
9          10       516.221073  
10          11       481.714845  
11          12       442.195541  
12          13       412.372292  
13          14       381.005891
```

```
[144]: # Elbow plot

plt.figure(figsize=(12,6))
plt.plot( clusters_df.num_clusters, clusters_df.cluster_errors, marker = "o" )
```

```
[144]: [<matplotlib.lines.Line2D at 0x1639d015da0>]
```



here we are taking k value as 4

```
[145]: kmeans = KMeans(n_clusters=3, n_init = 15, random_state=2345)
```

```
[147]: kmeans.fit(df7)
```

```
[147]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
              n_clusters=3, n_init=15, n_jobs=None, precompute_distances='auto',
              random_state=2345, tol=0.0001, verbose=0)
```

```
[148]: centroids = kmeans.cluster_centers_
centroids
```

```
[148]: array([[ -0.56000264,  0.91708692],
              [-0.26415144, -0.57202702],
              [ 1.22728261, -0.57114987]])
```

```
[150]: centroid_df = pd.DataFrame(centroids, columns = list(df7) )
centroid_df
```

```
[150]:      density      pH
0 -0.560003  0.917087
```

```
1 -0.264151 -0.572027
2  1.227283 -0.571150
```

```
[151]: ## creating a new dataframe only for labels and converting it into categorical
        ↪variable
df_labels = pd.DataFrame(kmeans.labels_ , columns = list(['labels']))

df_labels['labels'] = df_labels['labels'].astype('category')
```

```
[152]: # Joining the label dataframe with the Wine data frame to create
        ↪wine_df_labeled. Note: it could be appended to original dataframe
snail_df_labeled = df2.join(df_labels)
```

```
[153]: df_analysis = (snail_df_labeled.groupby(['labels'] , axis=0)).head(4177) # the
        ↪groupby creates a groupeddataframe that needs
        # to be converted back to dataframe. I am using .head(30000) for that
df_analysis
```

```
[153]:      labels
0         0
1         1
2         1
3         2
4         0
...      ...
1594      0
1595      0
1596      0
1597      0
1598      0

[1599 rows x 1 columns]
```

```
[154]: snail_df_labeled['labels'].value_counts()
```

```
[154]: 0    613
      1    581
      2    405
      Name: labels, dtype: int64
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```